Stats 232A project 1: Statistics of Natural Images

Due Date: Jan 21st, 11:59pm on CCLE

   This is a small project aiming at verifying some properties of natural image statistics that we discussed in Chapter 2. You can use any programming languages.

   Attached are four natural images (two from urban and two from countryside scenes).

   What to hand in: Submit a report on the CCLE site. Your score will be based on the quality of your results and the analysis (diagnostics of issues and comparisons) of these results in your report. Please also submit a zip file including your code, and make sure your code is consistent with results in report.

Problem 1 (High kurtosis and scale invariance, 3 points). For computational considerations, first convert image to grey level and re-scale the intensity to [0,31] i.e. 32 grey levels. Convolve the images with a gradient filter $\nabla_x I$, i.e. the intensity difference between two adjacent (horizontally or vertically). You can either pick any single image for the following steps or accumulate the histograms (average) over the 4 images.

1. Plot the histogram $H(z)$ for the difference against the horizontal axis $z \in [-31, +31]$. Then do a log-plot $\log H(z)$. [Some bins will be zero, you can assign ϵ for such bins in the log-plot].

2. Compute the mean, variance, and kurtosis for this histogram [Report the numeric numbers in your report].

3. Fit this histogram to a Generalized Gaussian distribution $\bar{e}^{|z/\sigma|^{\gamma}}$ and plot the fitted-curves super-imposed against the histogram. What is the value of γ in the fitted generalized Gaussian?

4. Plot the Gaussian distribution using the mean and the variance above, and super-impose this plot with the plots in step (1) above (i.e. plot the Gaussian and its log plot, this is easy to do in matlab).

5. Down-sample your image(s) by a 2×2 average (or simply sub-sample) the image. Plot the histogram and log histogram, and impose with the plots in step 1, to compare the difference. Repeat this down-sampling process 2-3 times.

6. Synthesize a uniform noise image, i.e. each pixel is drawn independently from a uniform number in $[0, 31]$. Repeat 1-2-4 to compare the difference between a noise image and a natural image. Do a 2×2 average instead of sub-sampling.

Problem 2. (Verify the 1/f power law observation in natural images, 3 points).

Do an FFT (Fast Fourier Transform) on the grey image I which returns a Fourier image $\hat{}(\xi, \eta)$ which is complex number matrix indexed by (ξ, η) for its horizontal and vertical frequencies. Compute the amplitude (modulus) of each complex number $A(\xi, \eta)$. Denote the frequency $f = \sqrt{\xi^2 + \eta^2}$, and transfer to a polar coordinate, and we calculate the total Fourier power $A^2(f)$ for each frequency (i.e. you need to discretize f, and calculate the $A^2(f)$ averaged over the ring for each f, stop f when the circle hits the boundary of the Fourier image).

1. Plot $\log A(f)$ against $\log f$. This should be close to a straight line for each image. Plot the curves for the 4 images in one figure for comparison.

2. Compute the integration (summation in discrete case) of $S(f_0) = \int_\Omega A^2(\xi, \eta) d\xi d\eta$ over the domain
$$\Omega(f_0) = \{(\xi, \eta) : f_0 \leq \sqrt{\xi^2 + \eta^2} \leq 2f_0\}$$
Plot $S(f_0)$ over f_0, the plot should fit to a horizontal line (with fluctuation) as $S(f_0)$ is supposed to be a constant over f_0.

Problem 3. (A 2D scale invariant world, 3 points). Suppose we simulate a toy 2D world where the images consist of only 1D line segments. In an image, a line segment is represented by its center (x_i, y_i), orientation θ_i and length r_i. The line segments are independently distributed with uniform probability for their centers and orientations. The length follows a probability $p(r) \propto 1/r^3$, i.e. a cubic power law. You may control the density of line by a Poisson distribution. That is, in each unit area, the number of line segments has a certain constant mean. [Hint: How to sample r from $p(r)$? Calculate the Cumulative Distribution function of $p(r)$, then draw a random number in [0,1].]

1. Simulate 1 image I_1 of size 1024×1024 pixels with a total N lines. (You need to record all the N lines whose centers are within a range of the image, truncate long lines and hide (discard) lines shorter than a pixel.)

2. Simulate 2 new images I_2 and I_3 of size 512×512 and 256×256 pixels respectively. I_2 and I_3 are down-sampled version of I_1 and are generated by shortening the N line segments in I_1 by 50% and 25% respectively (discard lines shorter than 1).

3. Crop 2 image patches of size 128×128 pixels randomly from each of the three images I_1, I_2, I_3 respectively. Plot these six images [draw the line segments in black on white background].

If you did it right [Please try !], the 6 images must look the same (i.e. you should not be able to tell what scale the 6 images are cropped from). So this 2D world is scale-invariant.

2